RL_filtered sharpens: late training proves known theorems more often but fewer theorems
pass@1 rises (sign test p = 0.0026)
n = 62 theorems
Mean pass@1 per theorem (%)
100
80
60
40
20
0
44.6%
58.3%
Early half
Late half
Coverage falls (McNemar p = 0.0063)
n = 62 theorems; 11 lost, 1 gained
Theorems proved at least once (count)
60
50
40
30
20
10
0
50
40
Early half
Late half